

Facial Expression Recognition for Cloud Robotics

Charlie Maclean
University of Cambridge
Cambridge, United Kingdom
cm927@cam.ac.uk

ABSTRACT

As social robots become widespread, it will be essential that they can classify the emotions of humans around them, in order to interact in a meaningful and helpful way. But limited hardware means they may have to offload video data to the cloud (cloud robotics), reducing the resolution of the content. This work focusses on evaluating the fitness of neural network models for cloud computing, specifically the effect of changing spatial and temporal resolution on the classification of emotion in video. We build different models to assess each model's performance when given lower resolution video. The results show that by applying a CNN-LSTM model to a TODO 7-class problem, we can achieve 73% accuracy on high resolution video, and maintain 66% accuracy with the lowest resolutions. To the best of our knowledge this is the first work that investigates the effect of changing both spatial and temporal resolution on video-based sentiment classification.

KEYWORDS

Affective computing, robotics, cloud computing, emotions, arousal, valence, resolution.

ACM Reference Format:

Charlie Maclean. 2021. Facial Expression Recognition for Cloud Robotics. In *Woodstock '18: ACM Symposium on Neural Gaze Detection, June 03–05, 2018, Woodstock, NY*. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/1122445.1122456>

1 INTRODUCTION

Social robots are becoming increasingly widespread, with uses in a wide range of locations, providing help in hospitals [5], care homes [7] and schools [24]. These robots are frequently required to interact meaningfully with humans, and in order to do so it is essential that they are able to classify emotions to react accordingly. However, many social robotic platforms lack the computational hardware required to perform classification [3]. Hence, it will likely become necessary to move to a cloud robotic framework, where sensing data is offloaded to the cloud and processed there. Unreliable network conditions mean we must be prepared for video data to enter the cloud at reduced spatial and temporal resolutions.

In past years, neural networks have become ubiquitous for classifying emotion, due to their ability to learn patterns humans would




Figure 1: Pepper by SoftBank Robotics, an example of a social robotic platform which has been trialled in care homes in the UK. Photograph by Dick Thomas, via creative commons. (source)

be unable to program in. However they can suffer from not being generalizable, especially if a network is trained in one domain, then deployed in another. For example, a network trained on high resolution data would be less effective at classifying low resolution data.

For classifying images, there is a large volume of work looking at using convolutional neural networks (CNNs). When applied to an image, a CNN convolves a filter with the pixel data, generating meaningful features. CNNs have found widespread use across various domains, including facial recognition [19] and object detection and classification [32]. Several architectures have been proposed offering impressive ability to learn features from video using purposes, for example the VGG16 network [30] and the ResNet50 network [12], both of which were able to achieve winning results in the ImageNet object detection and classification challenge [28].

For classifying videos, it is often vital to take temporal data into account, and as a result Recurrent Neural Networks (RNNs) [27] are a good choice. RNNs have some internal state, or memory, which they use to process sequences, learning patterns that may vary over time. A very popular architecture is Long Short Term Memory [13] (LSTM), which make use of gates to control the flow into and out of cells in the architecture. LSTMs have found wide usage, across speech recognition [11], market prediction [29] and handwriting recognition [10].

In this work we create a couple of classifiers which are tested on video at a variety of resolutions and frame-rates, in order to deduce which classifiers may be useful for cloud robotics. The classifiers

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Woodstock '18, June 03–05, 2018, Woodstock, NY

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/10.1145/1122445.1122456>

are tested on a 7-class video dataset, and our results show that we can achieve an accuracy of X on

The rest of the paper is structured as follows: Section 2 gives an overview of previous work on similar problems, Section 3 gives detail about the methodology employed in the study. Then, Section 4 discusses the results before Section 5 goes over future research directions. Finally, Section 6 concludes the paper with an overview of the findings.

2 RELATED WORK

In this section I give an overview of facial expression recognition techniques, first for images, and then for videos, followed by a section on classifying data at reduced resolutions.

2.1 Facial Expression Recognition in Images

Facial expressions have allowed humans to communicate their emotions amongst each other for years. There has been a large volume of research into the mechanisms which allow this non-verbal communication to happen. An early work by P. Ekman and W.V. Friesen introduced the Facial Action Coding System (FACS) [8], which described a list of facial action units - regions which change as a person changes their expression. The work further describes how a given facial expression could be described as a combination of action units. Following this work, P. Ekman et al. detailed how a mapping can be made between the facial action units and a person's emotions [6]. I will now detail the work that has been put into using machines to detect emotions from photos, splitting the research into those that make use of deep networks, and those before deep networks became widespread.

Before the use of deep networks, there were two main approaches to sentiment classification for images - rule-based methods, and appearance based methods. First, the rule-based methods were centred around detecting facial action units individually, and then piecing together the results from the facial action unit recognizers to derive an overall emotion, for example in Y. Tian et al.'s work [33]. These techniques suffered as recognizing an individual action unit is not an easy task. In the alternative approach, appearance based methods, some features are extracted from the overall face, and then those features are passed through a machine learning classifier. For example, M.S. Bartlett et al. found that they could achieve good results by extracting features using Gabor filters, followed by a Support Vector Machine (SVM) [2]. Additionally, J. Whitehall and C. Omlin showed that comparable results could be obtained in significantly quicker time by instead using Haar wavelets to obtain features before passing through a SVM [36].

In previous techniques, processing of emotions had been split into learning features, selecting features and then using a classifier to learn the patterns. The downside of taking that approach is that the first layers do not get feedback from the latter layers. Deep learning aims to solve that, by integrating the feature finding, selection and classification into one deep network that can be trained at once. One of the early papers making use of this technique was by P. Liu et al. [22], who suggested using a Boosted Deep Belief Network. This network consisted of several deep networks learning features, and some of these networks get boosted based on their

performance. Finally, in [20], Liu et al. introduced CNNs to the problem, with their CNN Ensemble network. The network consisted of three different convolutional networks, which proved to achieve better results than a single CNN.

2.2 Facial Expression Recognition in Videos

When classifying emotion in videos, there is additional information that can be extracted by accounting for the way the face changes over time. There are several papers which attempt to do this, which I will now go over, beginning with aggregation techniques. Aggregation techniques classify each frame within a video, and then combine the results with some sort of aggregator. In [14] and [15], S. Kahou et al. split the video into 10 sections, and within each section aggregate the frame predictions with an average. They go on to use a SVM to classify using all 10 aggregate predictions. An alternative method was introduced in [21] where M. Liu et al. showed that features from individual frames could be mapped to linear subspaces, covariance matrices or Gaussian distributions, allowing these to be passed to a support vector machine.

An alternative approach focussed on attempts to classify the level of emotional intensity present in an individual frame - the idea being that you could derive the emotion of a video base on the strongest emotions present. In [37] X. Zhao et al. propose a network which minimizes the differences between an emotion at low intensity and the same emotion at high intensity, as a way to get better classification of low intensity emotion. The downside of their technique was that it required a training set consisting of pairs of the same emotions at different intensities. To address this, in [4] J. Chen et al. used unsupervised clustering and a semisupervised SVM to detect peak and neutral frames in a large dataset.

Finally, deep spatio-temporal networks were introduced which use sequences of frames as inputs to the networks, in the hopes that we can learn more information from the temporal dynamics of the images. C3Ds [35] are the natural generalization of a CNN to 3 dimensions - instead of convolving an image with a 2D kernel, we convolve a sequence of frames with a 3D kernel. These 3D techniques have been brought to video emotion recognition, for example by X. Ouyang in [25]. An alternative approach was taken in [17] where D. Kim et al. tracked facial landmarks to generate trajectories which were used as features. Finally, a common approach is to use a CNN to learn spatial features of an image followed by a LSTM to learn the temporal features of the overall video, as in [16].

2.3 Reduced resolution classification

There have been a couple of studies looking at how to cope when the face that needs to be classified is at a low resolution. Firstly, Y. Tian [34] looked at how low spatial resolution would affect three steps in a facial expression analysis pipeline: face acquisition, facial data extraction and finally expression recognition. Also, R. Khan et al. produced a framework for expression recognition on low-resolution images [?], by suggesting a new set of features which work on low resolution images - called pyramid of local binary pattern. Finally, T. Vo et al. proposed the pyramid with super resolution architecture [?], which made use of super resolution networks in order to scale up low resolution images with minimal artefacts.

3 METHODOLOGY

In this section, the goal is to build a classifier which we can use to test video at multiple spatial and temporal resolutions. We first introduce the datasets used in the work, then discuss the image classifiers produced, before finally talking about the techniques used to classify the videos.

3.1 Datasets

For learning facial expressions of images I used the FER+ dataset [1] from E. Barsoum et al., a large collection of images first as part of the FER-2013 dataset [9]. The FER-2013 dataset consists of 36,685 48x48 greyscale images obtained by searching for different emotions on Google Images. The dataset is split into 7 categories - angry, sad, disgust, fear, happy, surprise and neutral. FER+ uses the same collection of images as in FER-2013, but updated the labels to be more accurate by crowd-sourcing the labelling and getting 10 taggers to label each image. In this work, we use majority voting to decide the ground truth of the images, discarding those for which there is no majority.

To learn facial expressions in videos, I used the RAVDESS [23] dataset by SR. Livingstone and FA. Russo. The dataset consists of 24 professional actors speaking two different statements while expressing one of eight emotions. The emotions include calm, happy, sad, angry, fearful, surprise, disgust and neutral. The data is provided at 720p, 30 fps with a blank background, in laboratory settings.

To pre-process the video data, we crop the images to a box around the face. In order to detect the bounding box around the face, we make use of Haar cascades. The purpose of this study is not to focus on how to detect faces at a reduced resolution as there is sufficient work already on that topic [38], [34].

3.2 Training Image Recognition

I tried two different CNN architectures, ResNet18 and ResNet50. They are 18 and 50 layer versions of the same family of Residual Networks, designed by K. He et al. [?]. They were designed to tackle image recognition tasks, and were significantly deeper than previously introduced architectures, which allowed them to gain accuracy in several recognition tasks. Further, despite the increase in depth, they managed to make the networks relatively easy to optimize.

Instead of training the network beginning with random parameters, I made use of transfer learning - where you use network weights obtained from a different task. This helps to bootstrap the network, giving it a good place to start learning the task you want. In our case we keep the base of the network the same and freeze it throughout training, so those weights remain unchanged. We use two different base models, which are publicly available - the first is trained on ImageNet [28], a massive dataset of over a million images, sorted into various nouns. Also, we use weights trained on the VGG Face dataset [26], which consists of 9,131 images of different people's faces. Finally, we add two dense layers to the head of the network which we do train.

To train the network we use stochastic gradient descent, and in order to obtain good hyperparameters we use Bayesian optimization [31]. Bayesian optimization represents the problem as a Gaussian process, which maps hyperparameters to the optimization

criteria. By doing this, it can focus the hyperparameter search on those regions which are most likely to lead to good results.

3.3 Training Video Recognition

To classify video, I decided to make use of the best network I previously used for classifying images with the logic that the patterns learnt there would be applicable in the video domain too. In order to extend the network that was designed for single images to video, with several frames, I made use of 3 different techniques. The first two are quite simplistic - I classify each frame of the video before then taking some aggregate of those results. For the first technique, this aggregate is a mean of all the classification scores. The second technique is where I take the max class for each frame, and then take the mode of these classifications.

The third technique I used was a LSTM network, appended to the output of the ResNet, to form a CNN-LSTM architecture. I used a LSTM with 1024 units, which means the LSTM outputs 1024 different features. I followed the LSTM with two dense layers for the final classification. The model was trained using the Adam optimizer [18], which promises to be well suited for optimization problems with a large number of parameters and data. As input to the network, I use 30 frames sequences from the video, and if the video was shorter than 40 frames I repeat frames until we get 30.

To preprocess the video, I created copies of the video at 3 different resolution and 3 different frame rates, for a total of 9 different clips for each input. The resolutions and frame rates were chosen such that they give a wide range of what may be possible over the internet. For resolutions, we had 720p (1280x720 px), 360p (480x360 px) and 144p (256x144 px), and for frame rates we had 30 fps, 15 fps and 5 fps. 720 fps at 30 fps is feasible with a fast internet connection, and low proximity to the server, and 144p at 5 fps would happen with a poor internet connection and/or large distance from a server.

4 EVALUATION

5 FUTURE WORK

6 CONCLUSION

REFERENCES

- [1] Emad Barsoum, Cha Zhang, Cristian Canton Ferrer, and Zhengyou Zhang. 2016. Training deep networks for facial expression recognition with crowd-sourced label distribution. In *Proceedings of the 18th ACM International Conference on Multimodal Interaction*. 279–283.
- [2] Marian Stewart Bartlett, Gwen Littlewort, Mark Frank, Claudia Lainscsek, Ian Fasel, and Javier Movellan. 2005. Recognizing facial expression: machine learning and application to spontaneous behavior. In *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*, Vol. 2. IEEE, 568–573.
- [3] Oya Celiktutan, Evangelos Sariyanidi, and Hatice Gunes. 2018. Computational Analysis of Affect, Personality, and Engagement in Human–Robot Interactions. In *Computer Vision for Assistive Healthcare*. Elsevier, 283–318.
- [4] Jingying Chen, Ruyi Xu, and Leyuan Liu. 2018. Deep peak-neutral difference feature for facial expression recognition. *Multimedia Tools and Applications* 77, 22 (2018), 29871–29887.
- [5] Diligent Robotics. 2019. Moxi – Diligent Robotics. <https://www.diligentrobots.com/moxi>
- [6] P Ekman, P.E.E.L. Rosenberg, P H D of Psychology Paul Ekman, E L Rosenberg, L.D.P.E.L. Rosenberg, and M B Smith. 1997. *What the Face Reveals: Basic and Applied Studies of Spontaneous Expression Using the Facial Action Coding System (FACS)*. Oxford University Press. <https://books.google.co.uk/books?id=fFGYs079-7YC>
- [7] ElliQ. 2019. ElliQ, the sidekick for happier aging – Intuition Robotics. <https://elliq.com/>

- [8] E Friesen and Paul Ekman. 1978. Facial action coding system: a technique for the measurement of facial movement. *Palo Alto* 3, 2 (1978), 5.
- [9] Ian J Goodfellow, Dumitru Erhan, Pierre Luc Carrier, Aaron Courville, Mehdi Mirza, Ben Hamner, Will Cukierski, Yichuan Tang, David Thaler, Dong-Hyun Lee, and Others. 2013. Challenges in representation learning: A report on three machine learning contests. In *International conference on neural information processing*. Springer, 117–124.
- [10] Alex Graves, Santiago Fernández, Marcus Liwicki, Horst Bunke, and Jürgen Schmidhuber. 2008. Unconstrained online handwriting recognition with recurrent neural networks. In *Advances in Neural Information Processing Systems 20, NIPS 2008*.
- [11] Alex Graves and Jürgen Schmidhuber. 2005. Framework phoneme classification with bidirectional LSTM and other neural network architectures. In *Neural Networks*, Vol. 18. Pergamon, 602–610. <https://doi.org/10.1016/j.neunet.2005.06.042>
- [12] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, Vol. 2016-Decem. IEEE Computer Society, 770–778. <https://doi.org/10.1109/CVPR.2016.90> arXiv:1512.03385
- [13] Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long short-term memory. *Neural computation* 9, 8 (1997), 1735–1780.
- [14] Samira Ebrahimi Kahou, Xavier Bouthillier, Pascal Lamblin, Caglar Gulcehre, Vincent Michalski, Kishore Konda, Sébastien Jean, Pierre Froumenty, Yann Dauphin, Nicolas Boulanger-Lewandowski, and Others. 2016. Emonets: Multimodal deep learning approaches for emotion recognition in video. *Journal on Multimodal User Interfaces* 10, 2 (2016), 99–111.
- [15] Samira Ebrahimi Kahou, Christopher Pal, Xavier Bouthillier, Pierre Froumenty, Caglar Gulcehre, Roland Memisevic, Pascal Vincent, Aaron Courville, Yoshua Bengio, Raul Chandias Ferrari, and Others. 2013. Combining modality specific deep neural networks for emotion recognition in video. In *Proceedings of the 15th ACM on International conference on multimodal interaction*. 543–550.
- [16] Dae Hoe Kim, Wissam J Baddar, Jinhyeok Jang, and Yong Man Ro. 2017. Multi-objective based spatio-temporal feature representation learning robust to expression intensity variations for facial expression recognition. *IEEE Transactions on Affective Computing* 10, 2 (2017), 223–236.
- [17] Dae Ha Kim, Min Kyu Lee, Dong Yoon Choi, and Byung Cheol Song. 2017. Multimodal emotion recognition using semi-supervised learning and multiple neural networks in the wild. In *Proceedings of the 19th ACM International Conference on Multimodal Interaction*. 529–535.
- [18] Diederik P Kingma and Jimmy Ba. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980* (2014).
- [19] Steve Lawrence, C Lee Giles, Ah Chung Tsoi, and Andrew D Back. 1997. Face recognition: A convolutional neural-network approach. *IEEE transactions on neural networks* 8, 1 (1997), 98–113.
- [20] Kuang Liu, Mingmin Zhang, and Zhigeng Pan. 2016. Facial expression recognition with CNN ensemble. In *2016 international conference on cyberworlds (CW)*. IEEE, 163–166.
- [21] Mengyi Liu, Ruiping Wang, Shaoxin Li, Shiguang Shan, Zhiwu Huang, and Xilin Chen. 2014. Combining multiple kernel methods on riemannian manifold for emotion recognition in the wild. In *Proceedings of the 16th International Conference on multimodal interaction*. 494–501.
- [22] Ping Liu, Shizhong Han, Zibo Meng, and Yan Tong. 2014. Facial expression recognition via a boosted deep belief network. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 1805–1812.
- [23] Steven R Livingstone and Frank A Russo. 2018. The Ryerson Audio-Visual Database of Emotional Speech and Song (RAVDESS): A dynamic, multimodal set of facial and vocal expressions in North American English. *PloS one* 13, 5 (2018), e0196391.
- [24] No Isolation. 2018. AV1 – the robot for children with long-term illness. <https://www.noisolation.com/global/av1> [https://www.noisolation.com/global/av1/](https://www.noisolation.com/global/av1/privacy-and-resources/%0Ahttps://www.noisolation.com/global/av1/)
- [25] Xi Ouyang, Shigenori Kawaai, Ester Gue Hua Goh, Shengmei Shen, Wan Ding, Huaiping Ming, and Dong-Yan Huang. 2017. Audio-visual emotion recognition using deep transfer learning and multiple temporal models. In *Proceedings of the 19th ACM international conference on multimodal interaction*. 577–582.
- [26] Omkar M Parkhi, Andrea Vedaldi, and Andrew Zisserman. 2015. Deep face recognition. (2015).
- [27] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. 1985. *Learning internal representations by error propagation*. Technical Report. California Univ San Diego La Jolla Inst for Cognitive Science.
- [28] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C Berg, and Li Fei-Fei. 2015. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)* 115, 3 (2015), 211–252. <https://doi.org/10.1007/s11263-015-0816-y>
- [29] Md. Saiful Islam and Emam Hossain. 2020. Foreign Exchange Currency Rate Prediction using a GRU-LSTM Hybrid Network. *Soft Computing Letters* (oct 2020), 100009. <https://doi.org/10.1016/j.socl.2020.100009>
- [30] Karen Simonyan and Andrew Zisserman. 2015. Very deep convolutional networks for large-scale image recognition. In *3rd International Conference on Learning Representations, ICLR 2015 - Conference Track Proceedings*. International Conference on Learning Representations, ICLR. arXiv:1409.1556 <http://www.robots.ox.ac.uk/>
- [31] Jasper Snoek, Hugo Larochelle, and Ryan P Adams. 2012. Practical bayesian optimization of machine learning algorithms. *arXiv preprint arXiv:1206.2944* (2012).
- [32] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, A Rabinovich, and Others. 2014. Going deeper with convolutions. *arXiv preprint arXiv:1409.4842* 1409 (2014).
- [33] Y-I Tian, Takeo Kanade, and Jeffrey F Cohn. 2001. Recognizing action units for facial expression analysis. *IEEE Transactions on pattern analysis and machine intelligence* 23, 2 (2001), 97–115.
- [34] Ying-li Tian. 2004. Evaluation of face resolution for expression analysis. In *2004 Conference on Computer Vision and Pattern Recognition Workshop*. IEEE, 82.
- [35] Du Tran, Lubomir Bourdev, Rob Fergus, Lorenzo Torresani, and Manohar Paluri. 2015. Learning spatiotemporal features with 3d convolutional networks. In *Proceedings of the IEEE international conference on computer vision*. 4489–4497.
- [36] Jacob Whitehill and Christian W Omlin. 2006. Haar features for FACS AU recognition. In *7th international conference on automatic face and gesture recognition (FG06)*. IEEE, 5–pp.
- [37] Xiangyun Zhao, Xiaodan Liang, Luoqi Liu, Teng Li, Yugang Han, Nuno Vasconcelos, and Shuicheng Yan. 2016. Peak-piloted deep network for facial expression recognition. In *European conference on computer vision*. Springer, 425–442.
- [38] Jun Zheng, Geovany A Ramirez, and Olac Fuentes. 2010. Face detection in low-resolution color images. In *International Conference Image Analysis and Recognition*. Springer, 454–463.