

# image

May 10, 2021

## 1 CNN Image Sentiment Analysis

Code to train a CNN image classifier.

Note that the final version of this code was run on the server as my computer was not powerful enough, and hence the latest version of my image classifier is found in `server_code/classify_image.py`

```
[1]: import matplotlib
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      from tensorflow.keras.layers import AveragePooling2D
      from tensorflow.keras.applications import ResNet50
      from tensorflow.keras.layers import Dropout
      from tensorflow.keras.layers import Flatten
      from tensorflow.keras.layers import Dense
      from tensorflow.keras.layers import Input
      from tensorflow.keras.models import Model
      from tensorflow.keras.models import load_model
      from tensorflow.keras.optimizers import SGD
      import tensorflow as tf
      from sklearn.preprocessing import LabelBinarizer
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report
      from imutils import paths
      import random
      import matplotlib.pyplot as plt
      import numpy as np
      import argparse
      import pickle
      import cv2
      import os
      import utils.data
      import utils.model
      from keras_vggface.vggface import VGGFace
      import gc
```

```
[ ]: import wandb
      from wandb.keras import WandbCallback
```

```
wandb.login()
```

Set out the parameters of this training run

```
[3]: # One of ravedess, ravedess-faces, fer+
dataset = "fer+"
FOUR_EMOTIONS = False

# One of RN50
EXPERIMENT = "RN50"

DIMS = (197,197)
```

We use transfer learning, on top of the ResNet CNN, using frames extracted from our videos to get the specific model.

```
[ ]: if EXPERIMENT in ['RN18-FER+', 'RN18-MS']:
    channels_first = True
else:
    channels_first = False

trainX, valX, testX, trainY, valY, testY, lb = utils.data.
    ↪load_img_dataset(dataset, channels_first, FOUR_EMOTIONS, dims=DIMS)

print(trainX.shape)

# Randomly change the train set so results are more generalizable
if EXPERIMENT in ['RN18-FER+', 'RN18-MS']:
    data_format = 'channels_first'
    train_augmentation = ImageDataGenerator(
        fill_mode="nearest",
        data_format=data_format)
else:
    data_format = 'channels_last'
    train_augmentation = ImageDataGenerator(
        rotation_range=30,
        zoom_range=0.15,
        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.15,
        horizontal_flip=True,
        fill_mode="nearest",
        data_format=data_format)

val_augmentation = ImageDataGenerator(data_format=data_format)
mean = np.array([123.68, 116.779, 103.939], dtype="float32")
train_augmentation.mean = mean
```

```
val_augmentation.mean = mean
```

```
[9]: def get_model():  
      return utils.model.get_model(EXPERIMENT, len(lb.classes_))
```

Set up the training procedure, using stochastic gradient descent. In the server version of this code we use a more complicated training procedure, where we first train the last few layers of the network, and then continue to train the whole network.

```
[10]: class MyCustomCallback(tf.keras.callbacks.Callback):  
      def on_epoch_end(self, epoch, logs=None):  
          gc.collect()  
  
      def train():  
          # default hyperparameters  
          config_defaults = {  
              'batch_size' : 32,  
              'learning_rate' : 0.0008475,  
              'epochs': 30,  
              'momentum' : 0.9,  
              'decay': 1e-4  
          }  
  
          wandb.init(project='sentiment', config=config_defaults)  
          config = wandb.config  
  
          config.architecture_name = EXPERIMENT  
          config.dataset = dataset  
  
          # Compile the model, using stochastic gradient descent optimization.  
          opt = SGD(lr=config.learning_rate, momentum=config.momentum, decay=config.  
→decay / config.epochs)  
          model = get_model()  
          model.compile(loss="categorical_crossentropy", optimizer=opt,  
                        metrics=["accuracy"])  
  
          # Now we can start training!  
          H = model.fit(  
              x = trainX,  
              y = trainY,  
              batch_size=config.batch_size,  
              steps_per_epoch = len(trainX) // config.batch_size,  
              validation_data = (valX, valY),  
              validation_steps = len(valX) // config.batch_size,  
              epochs = config.epochs,  
              callbacks = [WandbCallback()]  
          )
```

```
return model
```

Next, we setup a sweep of hyperparameters, using Bayesian optimization, implemented by weights and biases (wandb.ai).

```
[ ]: sweep_config = {
    "method": "bayes",
    "metric": {
        "name": "val_loss",
        "goal": "minimize"
    },
    "parameters":{
        "epochs": {
            "distribution": "int_uniform",
            "min": 20,
            "max": 40
        },
        "batch_size": {
            "distribution": "int_uniform",
            "min": 30,
            "max": 64
        },
        "learning_rate": {
            "distribution": "uniform",
            "min": 0.00001,
            "max": 0.001
        },
        "momentum": {
            "distribution": "uniform",
            "min": 0.9,
            "max": 0.99
        },
        "decay": {
            "distribution": "uniform",
            "min": 1e-6,
            "max": 1e-2
        }
    },
    "early_terminate" :
    {
        "type": "hyperband",
        "min_iter": 3
    }
}

wandb.sweep(sweep_config, project='sentiment')
```

```
wandb.agent(sweep_id, project='sentiment', function=train)
```